Connectionist Temporal Classification

Yi Wang

The Speech Recognition Problem

The input is an audio clip $\mathbf{x}$, as a sequence of frames:

$$\mathbf{x} = \{\mathbf{x}_1 \ldots \mathbf{x}_T\}$$

where each frame at time t, $\mathbf{x}_t$, consists of weights of G spectrograms:

$$\mathbf{x}_t = [x_{t,1} \ldots x_{t,G}]^T$$

We can design a RNN to output $\mathbf{y}_t$ for each input frame $\mathbf{x}_t$:

$$\mathbf{y} = \{\mathbf{y}_1 \ldots \mathbf{y}_T\}$$

where each output $\mathbf{y}_t$ is a probability distribution over an alphabet:

$$\mathbf{y}_t = [y_{t,1} \ldots y_{t,K}]^T$$

where $y_{t,k} = P(l_t = k)$, and l_t denotes the letter pronounced at time t.

Now, the problem comes – how can we interpret $\mathbf{y}$ as a sequence of letters.

At the first sight, this is straightforward – we just find a path

$$\boldsymbol{\pi} = \{\pi_1 \ldots \pi_T\}$$

where $\pi_t \in [1, K]$, by

$$\boldsymbol{\pi} = \arg\max_{\pi} P(\boldsymbol{\pi}|\mathbf{y})$$

where $P(\boldsymbol{\pi}|\mathbf{y})$ could be defined as, say,

$$P(\boldsymbol{\pi}|\mathbf{y}) = \prod_{t=1}^{T} y_{t,\pi_t}$$

or

$$P(\boldsymbol{\pi}|\mathbf{y}) = \prod_{t=1}^{T} y_{t,\pi_t} P(\pi_t|\pi_{t-1})$$

where $P(\pi_t|\pi_{t-1})$ comes from a pre-trained n-gram language model.

However, these are NOT what we need. Imagine that people are saying "a". There could be silence (or *blank*) before and/or after the pronunciation. So the path $\boldsymbol{\pi}$ could be

```
aaaaaaaa
aaaaa___
____aaaa
__aaaaa_
```

This does not even consider that the period of silences and pronunciations vary. So the following cases all correspond to "a":

```
aa_____
aaa_____
_aaa____
___aaaa_
_aaaaaaa
```

This tells that the output we really want is not $\mathbf{y}$ nor $\boldsymbol{\pi}$. Instead, it is some sequence called *label* and denoted by $\mathbf{l}$:

$$\mathbf{l} = \{l_1 \ldots l_S\}$$

where $S \leq T$.

For above example, all those paths correspond to the same label: $\mathbf{l} = \{a\}$.

Few more examples help reveal the relationship between path $\boldsymbol{\pi}$ and labels $\mathbf{l}$. All the following examples correspond to $\mathbf{l} = \{h, e\}$:

```
hhheeee
__heeee
_hee___
hh__eee
_hh_eee
h__ee__
__h_ee_
```

Please note that blanks could appear before, between, and after consequent appearances of "h" (and "e").

A little more complex example is that all the following paths correspond to label $\mathbf{l} = \{b, e, e\}$:

```
bbbeee_ee
_bb_ee__e
__bbbe_e_
```

except for

```
_b_eeeee
bbb__eeee
_bb_eee__
```

which correspond to $\mathbf{l} = \{b, e\}$. This example shows the importance of blank in separating consecutive letters.

In summary, the output of speech recognition is not alphabet probability distribution $\mathbf{y}$ or path $\boldsymbol{\pi}$, but label $\mathbf{l}$. And the CTC approach is all about infer $\mathbf{l}$ from $\mathbf{y}$, by integrating out $\boldsymbol{\pi}$:

$$P(\mathbf{l}|\mathbf{x}) = \sum_{\boldsymbol{\pi}} P(\mathbf{l}|\boldsymbol{\pi})P(\boldsymbol{\pi}|\mathbf{x})$$

where

$$P(\mathbf{l}|\boldsymbol{\pi}) = \begin{cases} 1 & \text{if } \boldsymbol{\pi} \text{ matches } \mathbf{l} \\ 0 & \text{otherwise} \end{cases}$$

$$P(\boldsymbol{\pi}|\mathbf{x}) = \prod_{t=1}^{T} y_{t,\pi_t}$$

Dynamic Time Warping

Because the length of $\mathbf{y}$ might differ from (often longer than) $\mathbf{l}$, so the inference of $\mathbf{l}$ from $\mathbf{y}$ is actually a dynamic time warping problem.

There is a dynamic programming algorithm that can solve this time warping problem. And this algorithm enables us to train the RNN using $\mathbf{x}$ and $\mathbf{l}$ pairs, instead of $\mathbf{x}$ and $\mathbf{y}$ pairs. This is very valuable because there is no need to segment $\mathbf{y}$ and align $\mathbf{y}$ with $\mathbf{x}$ before training.

When we say time warping, we mean that we want to map each frame $\mathbf{y}_t$ to some l_s. The second example above shows that we actually want to make $\mathbf{y}_t$ to either a l_s or a blank, because only if we do so, we have the chance to recognize successive letters like "ff" and "ee" in "coffee", and "ee" in "bee".

Here is a more detailed example about this. Suppose that we have an audio clip of "bee". The most probable path $\boldsymbol{\pi}$ is

```
___bbeeee____
```

we would like to wrap it to $\mathbf{l} = \{b, e, e\}$. It looks like we can map as

```
___bbeeee____
   \|||//
    |||
    bee
```

but the truth is that we do not have much information to prevent the algorithm from mapping all e's in $\boldsymbol{\pi}$ into the first e in $\mathbf{l}$:

```
___bbeeee____
   \||///
    ||
    bee
```

And another problem is that to whom should we map those blanks in $\boldsymbol{\pi}$?

A solution other than to map $\boldsymbol{\pi}$ to $\mathbf{l}$ is to map $\boldsymbol{\pi}$ to $\mathbf{l}'$, which is constructed by inserting blanks before, after, and into $\mathbf{l}$. For the above example, we have

$$\mathbf{l}' = \{, b, , e, , e, \}$$

and the warping could work as:

```
___bbeeee____
\\||/|||/|///
   || ||| /
  _b_e_e_
```

An a little more special case could be that $\boldsymbol{\pi}$ does not contain leading blanks. In this case, it is reasonable to map the first frame to "b", instead of the leading "_", of $\mathbf{l}'$:

```
 bbbeeee____
||/////////
||||| /
_b_e_e_
```

A similar edge case is that $\boldsymbol{\pi}$ does not have padding blanks, therefore no frame should be mapped to the padding blank of $\mathbf{l}'$. These two cases are what we need to care about in designing the algorithm that computes the map from $\boldsymbol{\pi}$ to $\mathbf{l}'$.

The Forward-Backward Algorithm

Here we derive the dynamic programming algorithm for computing $P(\mathbf{l}'|\mathbf{x})$, which, is notablly equivalent to $P(\mathbf{l}|\mathbf{x})$, because $\mathbf{l}'$ is constructed to be unique for the given $\mathbf{l}$.

The above time warping example shows that π_T could be mapped to either $l'_{|\mathbf{l}'|}$, the padding blank of $\mathbf{l}'$, or $l'_{|\mathbf{l}'|-1}$, which, according to the construction rule of $\mathbf{l}'$, is the last element of $\mathbf{l}$. So we have:

$$P(\mathbf{l}'|\mathbf{x}) = P(\mathbf{l}', \pi_T = l'_{|\mathbf{l}'|}|\mathbf{x}) + P(\mathbf{l}', \pi_T = l'_{|\mathbf{l}'|-1}|\mathbf{x})$$

We can generalize $P(\mathbf{l}', \pi_T = l'_{|\mathbf{l}'|}|\mathbf{x})$ and $P(\mathbf{l}', \pi_T = l'_{|\mathbf{l}'|-1}|\mathbf{x})$ to be

$$\alpha(t, s) = P(\mathbf{l}'_{1:s}, \pi_t = l'_s|\mathbf{x})$$

so

$$P(\mathbf{l}'|\mathbf{x}) = \alpha(T, |\mathbf{l}'|) + \alpha(T, |\mathbf{l}'| - 1)$$

Above time warping example also shows that π_1 could be mapped to either l'_1, the leading blank of $\mathbf{l}'$, or l'_2, the first element of $\mathbf{l}$. So we have

$$\alpha(1, 1) = y_{1,\_}$$
$$\alpha(1, 2) = y_{1,l_1}$$
$$\alpha(1, s) = 0, \forall s > 2$$

Here let us take an example. Suppose that $\mathbf{l}' = \{\_, h, \_, e, \_\}$ and a $\mathbf{y}$ (which, for the simplicity, is illustrated as a sequence of subscripts). It is reasonable to map π_1 to l'_1, the leading blank of $\mathbf{l}'$, if $\arg\max_k y_{1,k} = \text{``}\_\text{''}$:

```
123456789
|
|
_h_e_
```

or to $l'_2 = $ "h", if $\arg\max_k y_{1,k} = $ "h":

```
123456789
|
\
_h_e_
```

Then we think a step further – mapping π_2, or more generally, π_s. Roughly, it is reasonable to map π_t to

1. where π_{t-1} was mapped to, denoted by $l'_{s(t-1)}$,
2. the element next to $l'_{s(t-1)}$, denoted by $l'_{s(t-1)+1}$, or
3. $l'_{s(t-1)+2}$, if $l'_{s(t-1)+1} = $ "_"

Among these, case 3 is reasonable in case that we want to skip that blank $l'_{s(t-1)+1} = $ "_". An example is that when we want to map $\boldsymbol{\pi} = \{h, h, h, h, e, e, e\}$ to $\mathbf{l}' = \{\_, h, \_, e, \_\}$ – it is not mandatory to map any π_t to any blank in $\mathbf{l}'$, in order to recognize the word "he".

```
hhhheee
\|//|//
 |  /
_h_e_
```

But case 3 is not reasonable when $l'_{s(t-1)+2} = l'_{s(t-1)}$. In this case, we should not skip that blank. For example, to recognize the word "bee", we have $\mathbf{l}' = \{\_, b, \_, e, \_, e, \_\}$. In these case, if we skip over the blank between the two *e*s in $\mathbf{l}'$, we would misunderstand the double-e as a singel e. In the example below, even if frame $(\mathbf{y}_5)$ *sounds* more like *e* than blank, we want to map it to $l'_5 = $ "_", so to recognize the word "bee".

```
bbbeeeeee
\||||||///
 |/|||
_b_e_e_
```

Summarizing above three cases, we get the following generalization rule for $\alpha(t, s)$:

$$\alpha(t,s) = \begin{cases} y_{t,l'_s} \sum_{i=s-1}^{s} \alpha(t-1,i) & \text{if } l'_s = \_ \text{ or } l'_s = l'_{s-2} \\ y_{t,l'_s} \sum_{i=s-2}^{s} \alpha(t-1,i) & \text{otherwise} \end{cases}$$